

## Colab setup & dataset download (Roboflow → YOLOv8)

```

from IPython.display import clear_output
clear_output()
import os, json, shutil, random, glob
from pathlib import Path
import torch, os, glob, random, yaml, numpy as np
from IPython.display import display

# --- Clean install for NumPy 2.x compatible stack ---
%pip -q install -U pip setuptools wheel

# Remove conflicting builds
%pip -q uninstall -y numpy ultralytics onnx onnxruntime onnxruntime-gpu roboflow opencv-python

# Install NumPy 2 and libs that support it
%pip -q install "numpy==2.0.2" \
    "opencv-python-headless>=4.10.0.84" \
    "ultralytics>=8.3.184" \
    "onnx>=1.16.0" \
    "onnxruntime-gpu>=1.18.0" \
    "roboflow>=1.2.6" \
    ipywidgets>=8.1

#import IPython; IPython.get_ipython().kernel.do_shutdown(True) # will restart kernel

```

→ ━━━━━━━━━━━━━━━━ 1.8/1.8 MB 22.4 MB/s eta 0:00:00  
   ━ 1.2/1.2 MB 70.8 MB/s eta 0:00:00

ERROR: pip's dependency resolver does not currently take into account all the packages in file requirements.txt.  
 ipython 7.34.0 requires jedi>=0.16, which is not installed.

WARNING: Skipping ultralytics as it is not installed.

WARNING: Skipping onnx as it is not installed.

WARNING: Skipping onnxruntime as it is not installed.

WARNING: Skipping onnxruntime-gpu as it is not installed.

WARNING: Skipping roboflow as it is not installed.

```

import numpy, cv2, onnxruntime as ort, ultralytics, roboflow
print("numpy", numpy.__version__)
print("cv2", cv2.__version__)
print("onnxruntime providers:", ort.get_available_providers())
print("ultralytics", ultralytics.__version__)
import roboflow as rf; print("roboflow", rf.__version__)

```

→ Creating new Ultralytics Settings v0.0.6 file ✓  
 View Ultralytics Settings with 'yolo settings' or at '/root/.config/Ultralytics/settings.yaml'

```
Update Settings with 'yolo settings key=value', i.e. 'yolo settings runs_dir=path/to/dir
numpy 2.0.2
cv2 4.12.0
onnxruntime providers: ['TensorrtExecutionProvider', 'CUDAExecutionProvider', 'CPUExecutionProvider']
ultralytics 8.3.186
roboflow 1.2.6
```

## ▼ Gpu info

```
gpu_info = !nvidia-smi
gpu_info = '\n'.join(gpu_info)
if gpu_info.find('failed') >= 0:
    print('Not connected to a GPU')
else:
    print(gpu_info)
```

→ Mon Aug 25 17:53:16 2025

```
+-----
| NVIDIA-SMI 550.54.15                 Driver Version: 550.54.15      CUDA Version: 12.4
|-
| GPU  Name                  Persistence-M | Bus-Id      Disp.A  | Volatile Uncorr. EC
| Fan  Temp     Perf            Pwr:Usage/Cap |          Memory-Usage | GPU-Util  Compute M
|          |                               |              |           | GPU-Util  Compute M
|-----+
|  0  NVIDIA A100-SXM4-40GB        Off   | 00000000:00:04.0 Off |          0%       Default
| N/A   37C     P0            53W /  400W |          0MiB / 40960MiB |          |
|-----+
+-----+
| Processes:
| GPU  GI  CI          PID  Type  Process name             GPU Memori
| ID   ID
|-----+
| No running processes found
+-----+
```

## ▼ Roboflow: download the dataset you referenced (YOLOv8 format)

```
from roboflow import Roboflow
rf = Roboflow(api_key="ouWxulB15Sf3ls0jdSQD")
project = rf.workspace("roboflow-universe-projects").project("license-plate-recognition-rxg4")
version = project.version(11)
dataset = version.download("yolov8")
data_dir = Path(dataset.location)
data_dir
```

```
→ loading Roboflow workspace...
  loading Roboflow project...
  Downloading Dataset Version Zip in License-Plate-Recognition-11 to yolov8:: 100%|████████|
  Extracting Dataset Version Zip to License-Plate-Recognition-11 in yolov8:: 100%|████████|
  PosixPath('/content/License-Plate-Recognition-11')
```

## ▼ Inspect and show the dataset YAML path (Ultralytics uses this)

```
data_dir = Path("/content/License-Plate-Recognition-11")
yaml_path = next(data_dir.glob("*.yaml"))
print("Dataset YAML ->", yaml_path)
display(yaml.safe_load(open(yaml_path)))
```

```
→ Dataset YAML -> /content/License-Plate-Recognition-11/data.yaml
{'names': ['License_Plate'],
 'nc': 1,
 'roboflow': {'license': 'CC BY 4.0',
   'project': 'license-plate-recognition-rxg4e',
   'url': 'https://universe.roboflow.com/roboflow-universe-projects/license-plate-recognition-rxg4e/dataset/11'},
   'version': 11,
   'workspace': 'roboflow-universe-projects'},
   'test': '../test/images',
   'train': '../train/images',
   'val': '../valid/images'}
```

Dați dublu clic (sau apăsați pe Enter) pentru a edita

```
from ultralytics import YOLO
import torch

%pip -q install -U wandb
import os, wandb
os.environ.pop("WANDB_DISABLED", None)      # ensure not disabled
os.environ.pop("WANDB_MODE", None)           # ensure not offline/dryrun
wandb.login(relogin=True)                    # will prompt or reuse your key

from ultralytics import settings as yolo_settings
yolo_settings.update({'wandb': True})

MODEL_VARIANT = "n"                         # try 's' after this works (more accurate, slower)
IMGSZ = 416
EPOCHS = 50
BATCH = 16 if torch.cuda.is_available() else 8

model = YOLO(f"yolov8{MODEL_VARIANT}.pt")
results = model.train()
```

```
data=str(yaml_path),
imgsz=IMGSZ,
epochs=EPOCHS,
batch=BATCH,
patience=20,
device=0 if torch.cuda.is_available() else "cpu",
optimizer="SGD",
lr0=0.01,
weight_decay=5e-4,
hsv_h=0.015, hsv_s=0.7, hsv_v=0.4,
mosaic=1.0
)

# Optional: print the URL of the W&B run immediately
if wandb.run:
    print("W&B run:", wandb.run.get_url())
```

```
→ invalid escape sequence '\'
wandb: Logging into wandb.ai. (Learn how to deploy a W&B server locally: https://wandb.com)
wandb: You can find your API key in your browser here: https://wandb.ai/authorize
wandb: Paste an API key from your profile and hit enter: .....
wandb: WARNING If you're specifying your api key in code, ensure this code is not shar
wandb: WARNING Consider setting the WANDB_API_KEY environment variable, or running `wa
wandb: No netrc file found, creating one.
wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc
wandb: Currently logged in as: trinc-ciprian (trinc-ciprian-universitatea-de-vest-din)
Downloaded https://github.com/ultralytics/assets/releases/download/v8.3.0/yolov8n.pt
Ultralytics 8.3.186 🎉 Python-3.12.11 torch-2.8.0+cu126 CUDA:0 (NVIDIA A100-SXM4-40GB,
engine/trainer: agnostic_nms=False, amp=True, augment=False, auto_augment=randaugment,
Downloaded https://ultralytics.com/assets/Arial.ttf to '/root/.config/Ultralytics/Ari
Tracking run with wandb version 0.21.1
Run data is saved locally in /content/wandb/run-20250825_175418-qjblmo46
Syncing run train to Weights & Biases \(docs\)
View project at https://wandb.ai/trinc-ciprian-universitatea-de-vest-din-timisoara/Ultralytics
View run at https://wandb.ai/trinc-ciprian-universitatea-de-vest-din-timisoara/Ultralytics/runs/qjblmo46
Overriding model.yaml nc=80 with nc=1
```

	from	n	params	module	arg
0		-1	1	464	ultralytics.nn.modules.conv.Conv
1		-1	1	4672	ultralytics.nn.modules.conv.Conv
2		-1	1	7360	ultralytics.nn.modules.block.C2f
3		-1	1	18560	ultralytics.nn.modules.conv.Conv
4		-1	2	49664	ultralytics.nn.modules.block.C2f
5		-1	1	73984	ultralytics.nn.modules.conv.Conv
6		-1	2	197632	ultralytics.nn.modules.block.C2f
7		-1	1	295424	ultralytics.nn.modules.conv.Conv
8		-1	1	460288	ultralytics.nn.modules.block.C2f
9		-1	1	164608	ultralytics.nn.modules.block.SPPF
10		-1	1	0	torch.nn.modules.upsampling.Upsample
11	[ -1, 6 ]	1		0	ultralytics.nn.modules.conv.Concat
12		-1	1	148224	ultralytics.nn.modules.block.C2f
13		-1	1	0	torch.nn.modules.upsampling.Upsample
14	[ -1, 4 ]	1		0	ultralytics.nn.modules.conv.Concat
15		-1	1	37248	ultralytics.nn.modules.block.C2f
16		-1	1	36992	ultralytics.nn.modules.conv.Conv
17	[ -1, 12 ]	1		0	ultralytics.nn.modules.conv.Concat
18		-1	1	123648	ultralytics.nn.modules.block.C2f
19		-1	1	147712	ultralytics.nn.modules.conv.Conv
20	[ -1, 9 ]	1		0	ultralytics.nn.modules.conv.Concat
21		-1	1	493056	ultralytics.nn.modules.block.C2f
22	[ 15, 18, 21 ]	1	751507	ultralytics.nn.modules.head.Detect	[ 1,

Model summary: 129 layers, 3,011,043 parameters, 3,011,027 gradients, 8.2 GFLOPs

Transferred 319/355 items from pretrained weights

Freezing layer 'model.22.dfl.conv.weight'

**AMP:** running Automatic Mixed Precision (AMP) checks...

Downloading <https://github.com/ultralytics/assets/releases/download/v8.3.0/yolo11n.pt>

**AMP:** checks passed ✓

**train:** Fast image access ✓ (ping: 0.0±0.0 ms, read: 753.8±135.4 MB/s, size: 24.9 KB)

**train:** Scanning /content/License-Plate-Recognition-11/train/labels... 7057 image

**train:** New cache created: /content/License-Plate-Recognition-11/train/labels.ca

**albumentations:** Blur(p=0.01, blur\_limit=(3, 7)), MedianBlur(p=0.01, blur\_limit=(3, 7))

**val:** Fast image access ✓ (ping: 0.0±0.0 ms, read: 374.0±155.0 MB/s, size: 20.8 KB)

**val:** Scanning /content/License-Plate-Recognition-11/valid/labels... 2048 images, 3 bac  
**val:** New cache created: /content/License-Plate-Recognition-11/valid/labels.cache  
 Plotting labels to runs/detect/train/labels.jpg...  
**optimizer:** SGD(lr=0.01, momentum=0.937) with parameter groups 57 weight(decay=0.0), 64  
 Image sizes 416 train, 416 val  
 Using 8 dataloader workers  
 Logging results to **runs/detect/train**  
 Starting training for 50 epochs...

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
1/50	0.93G	1.289	1.371	1.222	3	416: 100% —
	Class all	Images 2048	Instances 2195	Box(P 0.936	R 0.888	mAP50 0.917 mAP50-9 0.5
2/50	1.14G	1.264	0.8446	1.183	2	416: 100% —
	Class all	Images 2048	Instances 2195	Box(P 0.936	R 0.886	mAP50 0.906 mAP50-9 0.5
3/50	1.15G	1.268	0.8345	1.188	1	416: 100% —
	Class all	Images 2048	Instances 2195	Box(P 0.892	R 0.839	mAP50 0.856 mAP50-9 0.5
4/50	1.17G	1.275	0.7852	1.203	2	416: 100% —
	Class all	Images 2048	Instances 2195	Box(P 0.951	R 0.875	mAP50 0.922 mAP50-9 0.5
5/50	1.18G	1.26	0.7278	1.204	1	416: 100% —
	Class all	Images 2048	Instances 2195	Box(P 0.965	R 0.875	mAP50 0.928 mAP50-9 0.6
6/50	1.19G	1.228	0.6857	1.179	2	416: 100% —
	Class all	Images 2048	Instances 2195	Box(P 0.958	R 0.891	mAP50 0.935 mAP50-9 0.
7/50	1.2G	1.212	0.6556	1.169	1	416: 100% —
	Class all	Images 2048	Instances 2195	Box(P 0.962	R 0.893	mAP50 0.928 mAP50-9 0.6
8/50	1.21G	1.201	0.641	1.172	1	416: 100% —
	Class all	Images 2048	Instances 2195	Box(P 0.959	R 0.879	mAP50 0.918 mAP50-9 0.6
9/50	1.22G	1.184	0.6237	1.156	1	416: 100% —
	Class all	Images 2048	Instances 2195	Box(P 0.964	R 0.903	mAP50 0.941 mAP50-9 0.6
10/50	1.23G	1.173	0.6037	1.15	2	416: 100% —
	Class all	Images 2048	Instances 2195	Box(P 0.959	R 0.879	mAP50 0.918 mAP50-9 0.6

## Validate + quick visual checks

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
11/50	1.25G	1.169	0.5923	1.148	2	416: 100% —

```
# 1) Enable W&B in Ultralytics
from ultralytics import settings as yolo_settings
yolo_settings.update({'wandb': True})

# 2) (optional) login once per session
import wandb; wandb.login(relogin=True)

# 3) Log a fresh run by validating your best checkpoint
from ultralytics import YOLO
best_pt = "runs/detect/train/weights/best.pt" # adjust if different
model = YOLO(best_pt)
metrics = model.val(imgsz=640, name="val_wandb_run") # creates a W&B run
print(metrics.results_dict)

from IPython.display import Image, display
import os

run_dir = "runs/detect/train2" # adjust if needed
for f in ["results.png", "confusion_matrix.png", "PR_curve.png", "F1_curve.png"]:
    path = os.path.join(run_dir, f)
    if os.path.exists(path):
        display(Image(filename=path, width=800))

# metrics (mAP, precision, recall, etc.)
metrics = model.val(data=str(yaml_path), imgsz=IMGSZ)
print(metrics.results_dict)

# visualize predictions on a handful of val images
val_imgs = glob.glob(f"{data_dir}/valid/images/*")
sample = random.sample(val_imgs, k=min(8, len(val_imgs)))
preds = model.predict(sample, imgsz=IMGSZ, conf=0.25, save=True, project="pred_samples", name="Saved examples under: pred_samples/run")
```

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
19/50	1.34G	1.117	0.5366	1.121	3	416: 100% —
	Class	Images	Instances	Box(P)	R	mAP50
	all	2048	2195	0.975	0.921	0.955

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
20/50	1.35G	1.113	0.5274	1.111	2	416: 100% —
	Class	Images	Instances	Box(P)	R	mAP50
	all	2048	2195	0.981	0.917	0.956

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
21/50	1.36G	1.094	0.5189	1.109	0	416: 100% —
	Class	Images	Instances	Box(P)	R	mAP50
	all	2048	2195	0.976	0.922	0.958

→ wandb: Logging into wandb.ai. (Learn how to deploy a W&B server locally: <https://wandb.ai>)  
wandb: Epoch GPU mem box\_loss cls\_loss dfl\_loss Instances Size  
wandb: You can find your API key in your browser here: <https://wandb.ai/authorize>  
wandb: 22/50 1.37G 1.098 0.5182 1.108 3 416: 100% —  
wandb: Paste an API key from your profile and hit enter: .....  
wandb: Class Images Instances Box(P) R mAP50 mAP50-9  
wandb: WARNING If you're specifying your api key in code, ensure this code is not shared  
wandb: WARNING Consider setting the WANDB\_API\_KEY environment variable, or running wanc  
wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc  
wandb: Epoch GPU mem box\_loss cls\_loss dfl\_loss Instances Size  
Ultralytics 8.3.186 Python-3.12.11-torch-2.8.0+cu126 CUDA:0 (NVIDIA A100-SXM4-40GB, 23750 1.38G 1.092 0.5095 1.103 4 416: 100% —  
Model summary (fused): 72 layers, 3,005,843 parameters, 0 gradients, 8.1 GELLOPs  
val: Fast image access ✓ (ping: 0.0±0.0 ms, read: 700.5±181.7 MB/s, size: 19.8 KB)  
val: Scanning /content/License-Plate-Recognition-11/valid/labels.cache... 2048 images, 24/50 1.39G 1.088 0.5082 1.106 4 416: 100% 0.665  
Speed: 0.1ms preprocess, 1.0ms inference, 0.0ms loss, 0.9ms postprocess per image  
Results saved to runs/detect/val\_wandb\_run  
{'metrics/precision(B)': 0.9660934496390103, 'metrics/recall(B)': 0.9346166791420323, 'me  
Ultralytics 8.3.186 Python-3.12.11-torch-2.8.0+cu126 CUDA:0 (NVIDIA A100-SXM4-40GB, 25/50 1.41G 1.084 0.5036 1.099 4 416: 100% —  
val: Fast image access ✓ (ping: 0.0±0.0 ms, read: 526.3±157.6 MB/s, size: 29.2 KB)  
val: Scanning /content/License-Plate-Recognition-11/valid/labels.cache... 2048 images, 26/50 1.42G 1.078 0.497 1.101 4 416: 100% —  
{'metrics/precision(B)': 0.980928431033285, 'metrics/recall(B)': 0.938496531435079, 'me  
0: 416x416 1 License\_Plate, 7.9ms  
1: 416x416 1 License\_Plate, 7.9ms  
2: 416x416 1 License\_Plate, 7.9ms  
3: 416x416 1 License\_Plate, 7.9ms  
4: 416x416 1 License\_Plate, 7.9ms  
5: 416x416 1 License\_Plate, 7.9ms  
6: 416x416 1 License\_Plate, 7.9ms  
7: 416x416 1 License\_Plate, 7.9ms  
Speed: 1.1ms preprocess, 7.9ms inference, 1.1ms postprocess per image at shape (1, 3, 41  
Results saved to pred\_samples/run  
Saved examples under: pred\_samples/run  
Epoch GPU mem box\_loss cls\_loss dfl\_loss Instances Size  
29/50 1.45G 1.067 0.4874 1.085 2 416: 100% —  
Class Images Instances Box(P) R mAP50 mAP50-9  
all 2048 2195 0.976 0.934 0.961 0.6

## Plot metrics directly in Python

Epoch	GPU mem	box_loss	cls_loss	dfl_loss	Instances	Size
30/50	1.46G	1.058	0.4848	1.08	2	416: 100% —

```
import matplotlib.pyplot as plt
```

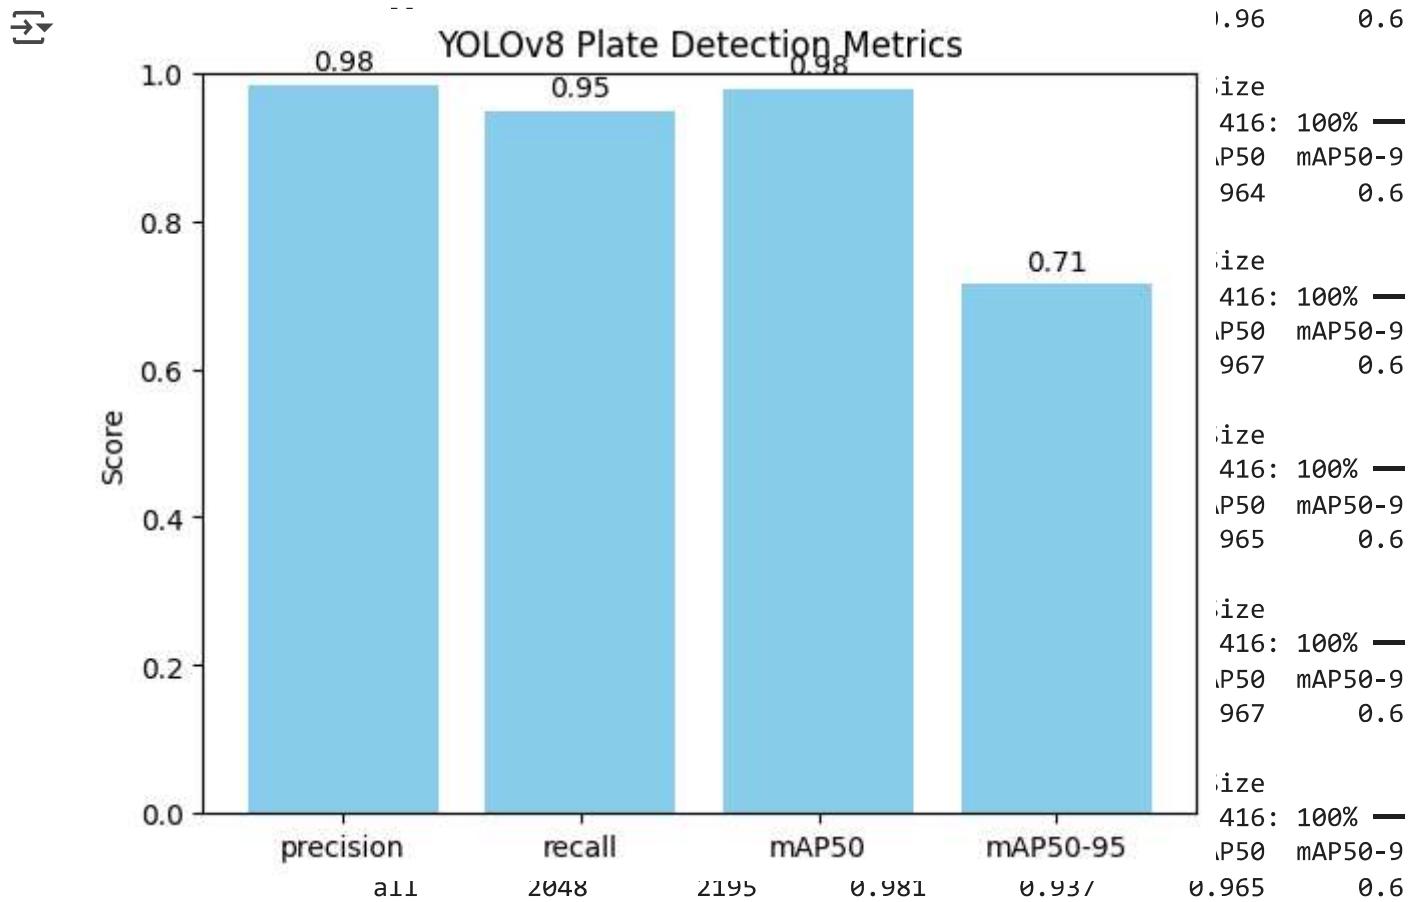
```
metrics = {
    'precision': 0.983,
    'recall': 0.949,
    'mAP50': 0.977,
    'mAP50-95': 0.714
}
```

```
plt.bar(metrics.keys(), metrics.values(), color="skyblue")
plt.ylim(0,1)
```

```

plt.ylabel("Score")
plt.title("YOLOv8 Plate Detection Metrics")
for k,v in metrics.items():
    plt.text(k, v+0.02, f"{v:.2f}", ha="center")
plt.show()

```



Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
1000	10G	1.011	0.4395	1.062	2	416: 100% —
	MS8G					mAP50 mAP50-9
	Class	Images	Instances	Box(P	R	
	all	2048	2195	0.973	0.939	0.964 0.6

Closing dataloader mosaic  
✓ Export NMS-free ONNX (needed for Hailo NMS later) [Implementation](#) [Example](#) [Run on Hailo](#) [Run on Hailo Dev Board](#) [Run on Hailo Dev Board \(NMS\)](#) [Blur limit=\(3, 7\)](#)

```

Epoch      GPU_mem      box_loss      cls_loss      dfl_loss      Instances      Size
best_pt = Path("runs/detect/train/weights/best.pt")
assert best_pt.exists(), "best.pt not found - check training run path."
onnx_path = YOLO(str(best_pt)).export(
    format="onnx",
    dynamic=False,          # static shapes
    opset=11,              # 12/13 OK
    imgsz=IMGSZ,
    half=False,             # keep FP32; Hailo will quantize
    simplify=True,
    nms=False               # <-- IMPORTANT: no builtin NMS
)
print("ONNX exported to:", onnx_path)

```

	Epoch	GPU mem	box_loss	cls_loss	dfl_loss	Instances	Size
→	Ultralytics 44/50	8.3.186 1.62G	Python-3.12.11 0.9857	Torch-2.8.0+cu126 0.3891	Fcu126 CPU (Intel Xeon 2.20GHz) 1.093	1 2195	416: 100% —
		💡 ProTip: Export to OpenVINO format for best performance on Intel hardware. Learn more	Class Images Instances Box(P R mAP50 mAP50-9				
	Model summary (fused)	all 72 layers	2048 3,005,843 parameters	2195 0.981	0 gradients, 8.1 GLOPs	0.938 0.964	0.6

**PyTorch:** starting from 'runs/detect/train/weights/best.pt' with input shape (1, 3, 416, 416)  
**requirements:** Ultralytics requirements: onnx>=1.17.0 & onnx>=1.18.0', onnxslim>=0.1.59, not 45/50

**requirements:** AutoUpdate success Class Images Instances Box(P R mAP50 mAP50-9  
**requirements:** AutoUpdate success all 2048 2195 0.982 0.934 0.965 0.6

WARNING ! **requirements:** Restart runtime or rerun command for updates to take effect

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
ONNX:	46/50	1.64G	0.9739	0.3818	1.092	1	416: 100% —
ONNX:	starting export with onnx 1.17.0 opset 11...	Class Images Instances Box(P R mAP50 mAP50-9					
ONNX:	slimming with onnxslim 0.1.65...	all 2048 2195 0.979 0.935 0.964 0.6					
ONNX:	export success ✓ 3.7s, saved as 'runs/detect/train/weights/best.onnx' (11.6 MB)						

Export complete (4-1s) Epoch 47/50 GPU\_mem box\_loss cls\_loss dfl\_loss Instances Size  
Results saved to content/runs/detect/train/weights Predict: yolo predict task=detect model=runs/detect/train/weights/best.onnx imgs Validate: yolo val task=detect model=runs/detect/train/weights/best.onnx imgs=41 Visualize: <https://netron.app>  
ONNX exported to: runs/detect/train/weights/best.onnx Epoch 48/50 GPU\_mem box\_loss cls\_loss dfl\_loss Instances Size  
416: 100% —

```
import onnx, onnxruntime as ort
m = onnx.load(onnx_path); onnx.checker.check_model(m)
sess = ort.InferenceSession(str(onnx_path), providers=["CUDAExecutionProvider", "CPUExecutionProvider"])
print("Input:", sess.get_inputs()[0].name, sess.get_inputs()[0].shape)
for o in sess.get_outputs():
    print("Output:", o.name, o.shape)
```

	Class	Images	Instances	Box(P	R	mAP50	mAP50-9
→	Input: images [1, 3, 416, 416]	2048	2195	0.984	0.937	0.966	0.7
	Output: output0 [1, 5, 3549]						

50 epochs completed in 0.536 hours.

Optimizer stripped from runs/detect/train/weights/last.pt, 6.2MB

## Build a calibration set (for Hailo quantization)

Validating runs/detect/train/weights/best.pt...

```
import cv2, os
calib_dir = Path("calib_npys"); calib_dir.mkdir(exist_ok=True)
```

```
def preprocess(path, size=IMGSZ):
    im = cv2.imread(path)[..., ::-1]      # BGR->RGB
    im = cv2.resize(im, (size, size))
    im = im.astype("float32")/255.0
    im = np.transpose(im, (2,0,1))        # HWC->CHW
    return im[None]                      # NCHW, batch=1
```

```
train_imgs = glob.glob(f"{data_dir}/train/images/*")
val_imgs   = glob.glob(f"{data_dir}/valid/images/*")
all_imgs = train_imgs + val_imgs
random.shuffle(all_imgs)
```

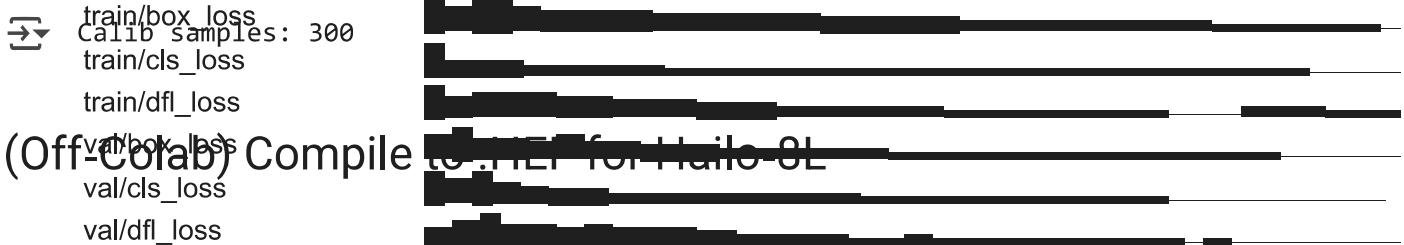
```

subset = all_imgs[:300] if len(all_imgs) >= 300 else all_imgs

for i, p in enumerate(subset):
    np.save(calib_dir/f"calib_{i:04d}.npy", preprocess(p))

print("Calib samples:", len(list(calib_dir.glob('*.*'))))

```



## Variables

Run summary:

```

MODEL_ONNX=/path/to/yolov8_plate.onnx CALIB_DIR=/path/to/calib_npys OUTDIR=/path/to/out
    lr/pg0          0.0003
mkdir $OUTDIR           0.0003
    lr/pg2          0.0003
    metrics/mAP50-95(B)  0.96597
    metrics/precision(B) 0.97924
    metrics/recall(B)   0.93895
hailom2 complete custom_yolov8_lpr
    model/GFLOPs        8.194
    model/parameters     3011043
-ckpt                         0.738
    model/speed_PyTorch(ms) 0.9440
$MODELONNX" --hw archhailo8l --batch-size1 --calib-path "
CALIB_DIR"
    train/cls_loss       0.35957
    train/dfl_loss        1.06887
-nms                          1.04907
    val/box_loss          0.39244
    val/cls_loss          1.11905
    val/dfl_loss
--output $OUTDIR/custom_yolov8_lpr.hef
-profile

```

View run **train** at: <https://wandb.ai/trinc-ciprian-universitatea-de-vest-din->

```

# --- Setup (run once) ---
# Point to the latest training run
run_dir = sorted(Path("runs/detect").glob("train"), key=lambda p: p.stat().st_mtime)[-1]
best_pt = run_dir / "weights" / "best.pt"
print("Using:", best_pt)

```

```

model = YOLO(str(best_pt))
IMGSZ = 640

```

→ Using: runs/detect/train/weights/best.pt

## ✓ 1) Test in Colab with your trained YOLO (image)

```
from pathlib import Path
import glob, random, IPython.display as disp# Option A: use a dataset image
img_path = random.choice(glob.glob("/content/License-Plate-Recognition-11/valid/images/*"))

# Option B: upload your own image
# from google.colab import files
# up = files.upload() # pick a local .jpg/.png
# img_path = list(up.keys())[0]

res = model.predict(img_path, imgsz=IMGSZ, conf=0.25, save=True, project="demo_out", name="j
print("Saved prediction to:", "demo_out/image")
disp.Image(filename=str(next((Path("demo_out/image")).glob("*.jpg")))))
```



image 1/1 /content/License-Plate-Recognition-11/valid/images/xemay351.jpg.rf.d165fe690c6  
 Speed: 1.8ms preprocess, 65.1ms inference, 1.6ms postprocess per image at shape (1, 3, 4)  
 Results saved to **demo\_out/image**  
 Saved prediction to: demo\_out/image



## ✓ 2) Test in Colab with your trained YOLO (image & video)

```
from ultralytics import YOLO
from pathlib import Path
import glob, random, cv2, numpy as np
import IPython.display as disp

# --- paths ---
DATA_DIR = Path("/content/License-Plate-Recognition-11")
IMG_DIR = DATA_DIR / "valid" / "images"      # you can switch to "test/images"
```

```
RUNS_DIR = Path("runs/detect")
OUT_DIR = Path("single_test_out"); OUT_DIR.mkdir(exist_ok=True, parents=True)

# --- load best checkpoint from the latest training run ---
run_dir = sorted(RUNS_DIR.glob("train"), key=lambda p: p.stat().st_mtime)[-1]
best_pt = run_dir / "weights" / "best.pt"
print("Using model:", best_pt)
model = YOLO(str(best_pt))
IMGSZ = 640

# --- pick ONE image ---
all_imgs = sorted(glob.glob(str(IMG_DIR/"*")))
assert all_imgs, f"No images found in {IMG_DIR}"
img_path = random.choice(all_imgs)
print("Testing on:", img_path)

# --- run detection & save Ultralytics-rendered result ---
res = model.predict(
    source=img_path,
    imgsz=IMGSZ,
    conf=0.25,
    save=True,
    project=str(OUT_DIR),
    name="pred"
)[0]

# the rendered image Ultralytics saved:
rendered = next((OUT_DIR / "pred").glob("*.jpg"))
print("Rendered with boxes ->", rendered)
disp.Image(filename=str(rendered), width=900)
```

→ Using model: runs/detect/train/weights/best.pt  
Testing on: /content/License-Plate-Recognition-11/valid/images/xemay2175.jpg.rf.f67a41bf  
  
image 1/1 /content/License-Plate-Recognition-11/valid/images/xemay2175.jpg.rf.f67a41bf95  
Speed: 1.9ms preprocess, 9.4ms inference, 2.1ms postprocess per image at shape (1, 3, 41)  
Results saved to **single\_test\_out/pred**  
Rendered with boxes -> single\_test\_out/pred/xemay2175.jpg.rf.f67a41bf991440fab0180fac4e6



### ▼ 3) Test in Colab with your trained YOLO (video)

```
# OCR step (host-side). Install once:  
%pip -q install easyocr > /dev/null  
  
import easyocr  
reader = easyocr.Reader(['en']) # add locale codes if needed, e.g., ['en','ro','de']  
  
# Load original image (BGR)  
img_bgr = cv2.imread(img_path)  
h, w = img_bgr.shape[:2]
```

```

# Copy for drawing OCR text
annot = img_bgr.copy()

# Iterate detections
texts = []
for b, conf, cls in zip(res.boxes.xyxy.cpu().numpy().astype(int),
                        res.boxes.conf.cpu().numpy(),
                        res.boxes.cls.cpu().numpy()):
    x1, y1, x2, y2 = b
    # clamp to image bounds
    x1, y1 = max(0,x1), max(0,y1)
    x2, y2 = min(w-1,x2), min(h-1,y2)
    crop = img_bgr[y1:y2, x1:x2]

    # OCR
    ocr_res = reader.readtext(crop)
    # pick the highest-score string if any
    if len(ocr_res):
        best = max(ocr_res, key=lambda t: t[2]) # (bbox, text, score)
        plate_text, score = best[1], best[2]
    else:
        plate_text, score = "", 0.0

    texts.append((b, plate_text, float(score), float(conf)))

# Draw box + OCR text
cv2.rectangle(annot, (x1,y1), (x2,y2), (0,255,0), 2)
label = f"{plate_text if plate_text else 'plate'} {conf:.2f} {score:.2f}"
cv2.putText(annot, label, (x1, max(0,y1-8)), cv2.FONT_HERSHEY_SIMPLEX, 0.6, (0,255,0), 2)

# Save and show
annot_path = OUT_DIR / "pred_with_ocr.jpg"
cv2.imwrite(str(annot_path), annot)
print("Annotated (with OCR) ->", annot_path)
disp.Image(filename=str(annot_path), width=900)

# Also print the raw detections + OCR strings
for b, plate_text, ocr_s, det_c in texts:
    print(f"Box {b} det_conf={det_c:.2f} OCR='{plate_text}' ocr_conf={ocr_s:.2f}")

```

→ WARNING:easyocr.easyocr:Downloading detection model, please wait. This may take several  
 Progress: |██████████| 100.0% CompleteWARNING:ea  
 Progress: |██████████| 100.0% CompleteAnnotated  
 Box [121 99 168 132] det\_conf=0.89 OCR='366,874' ocr\_conf=0.28

Dați dublu clic (sau apăsați pe Enter) pentru a edita

```
from pathlib import Path

# Look under runs/detect/
list(Path("runs/detect").glob("train*"))
```

→ [PosixPath('runs/detect/train')]

## locate

```
run_dir = sorted(Path("runs/detect").glob("train"), key=lambda p: p.stat().st_mtime)[-1]
print("Using run:", run_dir)
best_pt = run_dir / "weights" / "best.pt"
print("Best checkpoint exists:", best_pt.exists(), best_pt)
```

→ Using run: runs/detect/train  
Best checkpoint exists: True runs/detect/train/weights/best.pt

## Export to ONNX (no NMS)

```
from ultralytics import YOLO

model = YOLO(str(best_pt))
onnx_path = model.export(
    format="onnx",
    dynamic=False,
    opset=12,
    imgsz=640,
    half=False,
    simplify=True,
    nms=False      # important for Hailo
)

print("Exported ONNX path:", onnx_path)
```

→ Ultralytics 8.3.186 🚀 Python-3.12.11 torch-2.8.0+cu126 CPU (Intel Xeon 2.20GHz)  
Model summary (fused): 72 layers, 3,005,843 parameters, 0 gradients, 8.1 GFLOPs

**PyTorch:** starting from 'runs/detect/weights/best.pt' with input shape (1, 3, 640,

**ONNX:** starting export with onnx 1.17.0 opset 12...

**ONNX:** slimming with onnxslim 0.1.65...

**ONNX:** export success ✅ 1.6s, saved as 'runs/detect/weights/best.onnx' (11.3 MB)

Export complete (1.9s)

Results saved to /content/runs/detect/train/weights

```
Predict:      yolo predict task=detect model=runs/detect/train/weights/best.onnx imgs
Validate:     yolo val task=detect model=runs/detect/train/weights/best.onnx imgsz=64
Visualize:    https://netron.app
Exported ONNX path: runs/detect/train/weights/best.onnx
```

Dați dublu clic (sau apăsați pe Enter) pentru a edita

```
list((run_dir / "weights").glob("*.onnx"))
```

→ [PosixPath('runs/detect/train/weights/best.onnx')]

## ▼ Test the ONNX

```
import onnx, onnxruntime as ort, numpy as np

onnx_path = str(run_dir / "weights" / "best.onnx")
m = onnx.load(onnx_path); onnx.checker.check_model(m)

sess = ort.InferenceSession(onnx_path, providers=["CPUExecutionProvider"])
inp = sess.get_inputs()[0].name
print("Input:", inp, sess.get_inputs()[0].shape)

dummy = np.random.rand(1,3,640,640).astype("float32")
_ = sess.run(None, {inp: dummy})
print("ONNX forward pass OK")
```

→ Input: images [1, 3, 640, 640]  
ONNX forward pass OK

Dați dublu clic (sau apăsați pe Enter) pentru a edita

```
from pathlib import Path
import glob, random, cv2, numpy as np

DATA_DIR = Path("/content/License-Plate-Recognition-11")
IMGSZ = 640
CALIB_DIR = Path("calib_npys"); CALIB_DIR.mkdir(exist_ok=True, parents=True)

def preprocess(path, size=IMGSZ):
    im = cv2.imread(path)[..., ::-1]           # BGR->RGB
    im = cv2.resize(im, (size, size))
    im = (im.astype("float32")/255.0).transpose(2,0,1)  # CHW
    return im[None]                            # NCHW
```

```
train_imgs = glob.glob(str(DATA_DIR/"train/images/*"))
val_imgs   = glob.glob(str(DATA_DIR/"valid/images/*"))
all_imgs = train_imgs + val_imgs
random.shuffle(all_imgs)
subset = all_imgs[:300] if len(all_imgs) >= 300 else all_imgs

for i, p in enumerate(subset):
    np.save(CALIB_DIR/f"calib_{i:04d}.npy", preprocess(p))

len(list(CALIB_DIR.glob("*.npy")))
```

→ 300

Dați dublu clic (sau apăsați pe Enter) pentru a edita

```
from ultralytics import YOLO
from pathlib import Path
import glob, random, IPython.display as disp

best_pt = Path("runs/detect/train/weights/best.pt")
model = YOLO(str(best_pt))
IMGSZ = 640

img_path = random.choice(glob.glob("/content/License-Plate-Recognition-11/valid/images/*"))
print("Testing on:", img_path)

res = model.predict(source=img_path, imgsz=IMGSZ, conf=0.25, save=True,
                     project="demo_out", name="image")
print("Saved:", "demo_out/image")
disp.Image(filename=str(next((Path("demo_out/image")).glob("*.jpg"))), width=900)
```

→ Testing on: /content/License-Plate-Recognition-11/valid/images/xemay38.jpg.rf.8307d0b6b8f

image 1/1 /content/License-Plate-Recognition-11/valid/images/xemay38.jpg.rf.8307d0b6b881  
Speed: 1.7ms preprocess, 7.2ms inference, 1.6ms postprocess per image at shape (1, 3, 41)  
Results saved to demo\_out/image2  
Saved: demo\_out/image



## ▼ OCR detection

```
%pip -q install easyocr
import easyocr, cv2
reader = easyocr.Reader(['en']) # add your locale codes if needed

img = cv2.imread(img_path)
pred = model(img, imgsz=IMGSZ, conf=0.25)[0]
for box in pred.boxes.xyxy.cpu().numpy().astype(int):
    x1,y1,x2,y2 = box
    crop = img[y1:y2, x1:x2]
```

```
print("OCR:", reader.readtext(crop))
```

3

```
0: 416x640 1 License_Plate, 7.4ms  
Speed: 1.7ms preprocess, 7.4ms inference, 1.6ms postprocess per image at shape (1, 3, 41  
OCR: [[[np.int32(17), np.int32(7)], [np.int32(81), np.int32(7)], [np.int32(81), np.int32(7)]], [[np.int32(17), np.int32(7)], [np.int32(81), np.int32(7)], [np.int32(81), np.int32(7)]]]
```

```
!zip -r -q onnx_and_calib.zip runs/detect/train/weights/best.onnx calib_npys
```

Dați dublu clic (sau apăsați pe Enter) pentru a edita

```
from ultralytics import YOLO
import glob, random
from pathlib import Path
from IPython.display import Image

model = YOLO("runs/detect/train/weights/best.pt")
img = random.choice(glob.glob("/content/License-Plate-Recognition-11/valid/images/*"))
pred = model.predict(img, imgsz=640, conf=0.25, save=True, project="demo_out", name="img")
Image(filename=str(next((Path("demo_out/img")).glob("*.jpg"))), width=900)
```



image 1/1 /content/License-Plate-Recognition-11/valid/images/00cac7ea145fc734.jpg.rf.f26  
Speed: 2.5ms preprocess, 65.1ms inference, 2.0ms postprocess per image at shape (1, 3, 4)  
Results saved to **demo\_out/img**



## ▼ Example from Youtube

```
!pip -q install yt-dlp ultralytics opencv-python-headless
```

```
import yt_dlp
```

```
url = "https://www.youtube.com/watch?v=70IoL7irIF8" # traffic driving sample

ydl_opts = {
    "format": "best[ext=mp4]",
    "outtmpl": "input_video.%({ext}s"
}
with yt_dlp.YoutubeDL(ydl_opts) as ydl:
    ydl.download([url])
```

→ [youtube] Extracting URL: <https://www.youtube.com/watch?v=70IoL7irIF8>  
[youtube] 70IoL7irIF8: Downloading webpage  
[youtube] 70IoL7irIF8: Downloading tv client config  
[youtube] 70IoL7irIF8: Downloading player 6742b2b9-main  
[youtube] 70IoL7irIF8: Downloading tv player API JSON  
[youtube] 70IoL7irIF8: Downloading tv simply player API JSON  
[info] 70IoL7irIF8: Downloading 1 format(s): 18  
[download] Sleeping 2.00 seconds as required by the site...  
[download] Destination: input\_video.mp4  
[download] 100% of 3.79MiB in 00:00:00 at 11.96MiB/s

```
from ultralytics import YOLO

# load your trained model
model = YOLO("runs/detect/train/weights/best.pt")

# run detection on video
results = model.predict(
    source="input_video.mp4",    # downloaded video
    imgsz=640,
    conf=0.25,
    save=True,
    project="demo_out",
    name="youtube_video"
)

print("Processed video saved to:", "demo_out/youtube_video/input_video.mp4")
```

→ WARNING !  
inference results will accumulate in RAM unless `stream=True` is passed, causing potential errors for large sources or long-running streams and videos. See <https://docs.ultralytics.com>

Example:

```
results = model(source=..., stream=True) # generator of Results objects
for r in results:
    boxes = r.boxes # Boxes object for bbox outputs
    masks = r.masks # Masks object for segment masks outputs
    probs = r.probs # Class probabilities for classification outputs
```

```
video 1/1 (frame 1/1253) /content/input_video.mp4: 384x640 1 License_Plate, 62.2ms
video 1/1 (frame 2/1253) /content/input_video.mp4: 384x640 (no detections), 7.0ms
```

```
!ffmpeg -i /content/demo_out/youtube_video/input_video.avi \
    -vcodec libx264 -crf 23 -preset veryfast \
    -acodec aac -strict -2 \
    /content/demo_out/youtube_video/input_video.mp4 -y
```

```
→ ffmpeg version 4.4.2-0ubuntu0.22.04.1 Copyright (c) 2000-2021 the FFmpeg develop  
built with gcc 11 (Ubuntu 11.2.0-19ubuntu1)  
configuration: --prefix=/usr --extra-version=0ubuntu0.22.04.1 --toolchain=hardened --
```

```

libavutil      56. 70.100 / 56. 70.100
libavcodec     58.134.100 / 58.134.100
libavformat    58. 76.100 / 58. 76.100
libavdevice    58. 13.100 / 58. 13.100
libavfilter     7.110.100 / 7.110.100
libswscale      5.  9.100 / 5.  9.100
libswresample   3.  9.100 / 3.  9.100
libpostproc    55.  9.100 / 55.  9.100
Input #0, avi, from '/content/demo_out/youtube_video/input_video.avi':
  Metadata:
    software       : Lavf59.27.100
  Duration: 00:00:41.77, start: 0.000000, bitrate: 15446 kb/s
  Stream #0:0: Video: mjpeg (Baseline) (MJPEG / 0x47504A4D), yuvj420p(pc, bt470bg/unknowr
  Stream mapping:
    Stream #0:0 -> #0:0 (mjpeg (native) -> h264 (libx264))
  Press [q] to stop, [?] for help
[libx264 @ 0x5ae680c5cd40] using cpu capabilities: MMX2 SSE2Fast SSSE3 SSE4.2 AVX FMA3 E
[libx264 @ 0x5ae680c5cd40] profile High, level 3.0, 4:2:0, 8-bit
[libx264 @ 0x5ae680c5cd40] 264 - core 163 r3060 5db6aa6 - H.264/MPEG-4 AVC codec - Copy]
Output #0, mp4, to '/content/demo_out/youtube_video/input_video.mp4':
  Metadata:
    software       : Lavf59.27.100
    encoder        : Lavf58.76.100
  Stream #0:0: Video: h264 (avc1 / 0x31637661), yuvj420p(pc, bt470bg/unknown/unknown, pr
  Metadata:
    encoder        : Lavc58.134.100 libx264
  Side data:
    cpb: bitrate max/min/avg: 0/0/0 buffer size: 0 vbv_delay: N/A
frame= 1253 fps=526 q=-1.0 Lsize=    4802kB time=00:00:41.66 bitrate= 944.0kbits/s speec
video:4786kB audio:0kB subtitle:0kB other streams:0kB global headers:0kB muxing overhead
[libx264 @ 0x5ae680c5cd40] frame I:6      Avg QP:23.90  size: 49560
[libx264 @ 0x5ae680c5cd40] frame P:343    Avg QP:26.51  size: 10246
[libx264 @ 0x5ae680c5cd40] frame B:904    Avg QP:30.79  size: 1204
[libx264 @ 0x5ae680c5cd40] consecutive B-frames: 2.8% 2.4% 1.9% 92.9%
[libx264 @ 0x5ae680c5cd40] mb I  I16..4: 0.7% 20.5% 78.8%
[libx264 @ 0x5ae680c5cd40] mb P  I16..4: 0.4% 6.1% 0.8% P16..4: 48.8% 17.0% 13.3% E
[libx264 @ 0x5ae680c5cd40] mb B  I16..4: 0.1% 0.3% 0.0% B16..8: 12.5% 5.4% 1.0% C
[libx264 @ 0x5ae680c5cd40] 8x8 transform intra:71.5% inter:27.2%
[libx264 @ 0x5ae680c5cd40] coded y,uvDC,uvAC intra: 76.1% 64.5% 17.1% inter: 14.4% 9.0%
[libx264 @ 0x5ae680c5cd40] i16 v,h,dc,p: 21% 41% 34% 5%
[libx264 @ 0x5ae680c5cd40] i8 v,h,dc,ddl,ddr,vr,hd,vl,hu: 23% 34% 31% 2% 2% 1% 3% 1
[libx264 @ 0x5ae680c5cd40] i4 v,h,dc,ddl,ddr,vr,hd,vl,hu: 19% 28% 12% 5% 6% 6% 9% E
[libx264 @ 0x5ae680c5cd40] i8c dc,h,v,p: 45% 36% 15% 4%
[libx264 @ 0x5ae680c5cd40] Weighted P-Frames: Y:9.6% UV:0.9%
[libx264 @ 0x5ae680c5cd40] kb/s:938.60

```

```

from IPython.display import Video
Video("/content/demo_out/youtube_video/input_video.mp4", embed=True, width=800)

```



0:00 / 0:41



## ▼ Mount Drive & point to file

```
from google.colab import drive
drive.mount('/content/drive', force_remount=True)

# <<< path in Drive >>>
INPUT_MOV = "/content/drive/MyDrive/Video_H/IMG_0130.MOV"
```

→ Mounted at /content/drive

## ▼ convert MOV → MP4 for smoother processing

```
!ffmpeg -y -i "$INPUT_MOV" -vcodec libx264 -crf 20 -preset veryfast -acodec aac /content/inp
```

→ ffmpeg version 4.4.2-0ubuntu0.22.04.1 Copyright (c) 2000-2021 the FFmpeg developers  
built with gcc 11 (Ubuntu 11.2.0-19ubuntu1)  
configuration: --prefix=/usr --extra-version=0ubuntu0.22.04.1 --toolchain=hardened  
libavutil 56. 70.100 / 56. 70.100  
libavcodec 58.134.100 / 58.134.100  
libavformat 58. 76.100 / 58. 76.100

```

libavdevice      58. 13.100 / 58. 13.100
libavfilter       7.110.100 / 7.110.100
libswscale        5.  9.100 / 5.  9.100
libswresample     3.  9.100 / 3.  9.100
libpostproc       55. 9.100 / 55. 9.100
Input #0, mov,mp4,m4a,3gp,3g2,mj2, from '/content/drive/MyDrive/Video_H/IMG_0130.MOV':
  Metadata:
    major_brand     : qt
    minor_version   : 0
    compatible_brands: qt
    creation_time   : 2025-08-24T13:55:48.000000Z
    com.apple.quicktime.location.accuracy.horizontal: 7.931697
    com.apple.quicktime.location.ISO6709: +45.7326+021.2351+092.789/
    com.apple.quicktime.make: Apple
    com.apple.quicktime.model: iPhone 11
    com.apple.quicktime.software: 17.5.1
    com.apple.quicktime.creationdate: 2025-08-24T16:27:22+0300
    com.apple.photos.originating.signature: AXiC42ncYYXTiJFZFWIocCpovxdo
Duration: 00:00:09.77, start: 0.000000, bitrate: 15319 kb/s
Stream #0:0(und): Audio: aac (LC) (mp4a / 0x6134706D), 44100 Hz, stereo, fltp, 152 k
  Metadata:
    creation_time   : 2025-08-24T13:55:48.000000Z
    handler_name    : Core Media Audio
    vendor_id       : [0][0][0][0]
Stream #0:1(und): Video: h264 (High) (avc1 / 0x31637661), yuv420p(tv, bt709), 1920x1
  Metadata:
    creation_time   : 2025-08-24T13:55:48.000000Z
    handler_name    : Core Media Video
    vendor_id       : [0][0][0][0]
    encoder         : H.264
Stream #0:2(und): Data: none (mebx / 0x7862656D), 0 kb/s (default)
  Metadata:
    creation_time   : 2025-08-24T13:55:48.000000Z
    handler_name    : Core Media Metadata
Stream #0:3(und): Data: none (mebx / 0x7862656D), 0 kb/s (default)
  Metadata:
    creation_time   : 2025-08-24T13:55:48.000000Z
    handler_name    : Core Media Metadata
Stream #0:4(und): Data: none (mebx / 0x7862656D), 34 kb/s (default)
  Metadata:
    creation_time   : 2025-08-24T13:55:48.000000Z
    handler_name    : Core Media Metadata
Stream mapping:
  Stream #0:1 -> #0:0 (h264 (native) -> h264 (libx264))
  Stream #0:0 -> #0:1 (aac (native) -> aac (native))
Press [q] to stop, [?] for help
[libx264 @ 0x58f7840d1d40] using cpu capabilities: MMX2 SSE2Fast SSSE3 SSE4.2 AVX FMA3
[libx264 @ 0x58f7840d1d40] profile High, level 4.0, 4:2:0, 8-bit
[libx264 @ 0x58f7840d1d40] 264 - core 163 r3060 5db6aa6 - H.264/MPEG-4 AVC codec - Cop
Output #0, mp4, to '/content/input_videoP.mp4':
  Metadata:
    .
    .

```

## ▼ Run YOLO plate detector on the video

```
from ultralytics import YOLO
from pathlib import Path

model = YOLO("runs/detect/train/weights/best.pt") # your trained model

res = model.predict(
    source="/content/input_videoP.mp4",    # the mp4 we just made
    imgsz=640,
    conf=0.25,
    save=True,
    project="demo_out",
    name="drive_videoP",
    vid_stride=1
)

# Ultralytics writes output into demo_out/drive_video/
out_dir = Path("demo_out/drive_videoP")
print("Outputs in:", out_dir)
print("Files:", [p.name for p in out_dir.iterdir()])
```



#### WARNING !

inference results will accumulate in RAM unless `stream=True` is passed, causing potential errors for large sources or long-running streams and videos. See <https://docs.ultralytics.com>

#### Example:

```
results = model(source=..., stream=True) # generator of Results objects
for r in results:
    boxes = r.boxes # Boxes object for bbox outputs
    masks = r.masks # Masks object for segment masks outputs
    probs = r.probs # Class probabilities for classification outputs

video 1/1 (frame 1/293) /content/input_videoP.mp4: 384x640 4 License_Plates, 7.9ms
video 1/1 (frame 2/293) /content/input_videoP.mp4: 384x640 4 License_Plates, 10.7ms
video 1/1 (frame 3/293) /content/input_videoP.mp4: 384x640 3 License_Plates, 7.2ms
video 1/1 (frame 4/293) /content/input_videoP.mp4: 384x640 3 License_Plates, 7.0ms
video 1/1 (frame 5/293) /content/input_videoP.mp4: 384x640 3 License_Plates, 7.0ms
video 1/1 (frame 6/293) /content/input_videoP.mp4: 384x640 4 License_Plates, 7.0ms
video 1/1 (frame 7/293) /content/input_videoP.mp4: 384x640 4 License_Plates, 7.0ms
video 1/1 (frame 8/293) /content/input_videoP.mp4: 384x640 3 License_Plates, 7.1ms
video 1/1 (frame 9/293) /content/input_videoP.mp4: 384x640 3 License_Plates, 7.1ms
video 1/1 (frame 10/293) /content/input_videoP.mp4: 384x640 3 License_Plates, 10.5ms
video 1/1 (frame 11/293) /content/input_videoP.mp4: 384x640 3 License_Plates, 7.2ms
video 1/1 (frame 12/293) /content/input_videoP.mp4: 384x640 4 License_Plates, 7.0ms
video 1/1 (frame 13/293) /content/input_videoP.mp4: 384x640 4 License_Plates, 7.0ms
```